



Module	Data Science Project (Research)
Code	NIB7150CEM
Submission	Research Proposal
Student Name	Samanthilaka K P R
Student ID	COMScDS232P-014
Email	COMScDS232P-014@student.nibm.lk

MSc Data Science
Coventry University UK

Mitigating Hallucinations in Generative AI: through Agents, Retrieval-Augmented Generation (RAG) and support results with enhanced Referencing and Reasoning for High-Stakes Environments.

Introduction

Generative AI, powered by Large Language Models (LLMs), has transformed how industries handle tasks, processes, and decision-making with advancement of language understanding and generation. These technologies are now commonly used in customer service to education, where the speed and efficiency exceeds human capabilities. However, as the usage of LLMs grows, so do the challenges associated with its output. One of the most pressing issues is hallucination—a phenomenon where these models produce statements which appear to be confident and convincing but are inaccurate or fabricated. (Dhuliawala et al., 2023) Such outputs can severely weaken the trust in these models, especially in high-stakes environments.

Industries like healthcare, finance, and legal require outputs should be coherent, factually correct and reliable. Any form of hallucination in these industries can lead to serious consequences. This limitation of current generative AI systems reduces spread of AI system in any environment. While some progress has been made in refining LLMs to reduce hallucinations, still significant levels of hallucination can be visible in all major LLMs available. (Manakul et al., 2023)

Problem Statement

In high-stakes environments such as healthcare, finance, and legal, the accuracy and reliability of AI-driven language models are vital. However, large language models (LLMs) frequently produce "hallucinations" — confidently generated outputs that are factually incorrect or misleading. This poses significant risks in such contexts where incorrect information could lead to severe consequences. Existing mitigation techniques lack comprehensive frameworks to reduce these hallucinations, especially in situations requiring precise information and accountability.

This research aims to develop an integrated approach leveraging agents and Retrieval-Augmented Generation (RAG) techniques to minimize hallucinations in LLMs. The proposed solution will focus on refining real-time referencing and enhancing reasoning capabilities to ensure verifiable, contextually accurate outputs.

Objectives

Develop and evaluate techniques that help Large Language Models (LLMs) to generate more accurate, dependable and verifiable responses.

- Identify root causes and limitations for hallucinations
- Design and implement an agent system that manage and refine generation process
- Integrate RAG to reinforce LLM with external and factual data in response generation
- Identify and integrate referencing and reasoning mechanisms
- Evaluate effectiveness in high-stake environments for accuracy, reliability

Literature Review

The **concept of hallucination** traces its roots to the fields of pathology and psychology and is defined as *the perception of an entity or event that is absent in reality*. Generally, hallucinations in natural language generation tasks can be categorized into two primary types: *intrinsic hallucination* and *extrinsic hallucination*. Specifically, *intrinsic hallucinations* pertain to the outputs of LLMs that conflict with the source content. Conversely, *extrinsic hallucinations* refer to the LLM generations that cannot be verified from the source content. (Huang et al., 2023)

Factuality Hallucination. The emergence of LLMs marks a significant shift from traditional task-specific toolkits to AI assistants that have a heightened focus on open-domain interactions. (Huang et al., 2023)

- **Factual Inconsistency** refers to situations where the LLM's output contains facts that can be grounded in real-world information, but present contradictions.
- **Factual Fabrication** refers to instances where the LLM's output contains facts that are unverifiable against established real-world knowledge.

Faithfulness Hallucination. LLMs are inherently trained to align with user instructions. As the use of LLMs shifts towards more user-centric applications, ensuring their consistency with user provided instructions and contextual information becomes increasingly vital. (Huang et al., 2023)

- **Instruction inconsistency** refers to the LLM's outputs that deviate from a user's directive.
- **Context inconsistency** points to instances where the LLM's output is unfaithful with the user's provided contextual information.
- **Logical inconsistency** underscores when LLM outputs exhibit internal logical contradictions, often observed in reasoning tasks.

Retrieval Augmented Generation (RAG) has emerged as a promising strategy to mitigate hallucinations in LLMs. (Huang et al., 2023) While autonomous agents have long been considered a promising approach to achieving artificial general intelligence (AGI) (Han et al., 2024). Combining these two techniques together might end the issues of hallucination and achieve 100% accuracy in high-risk domains.

Retrieval-Augmented Generation (RAG) mitigates hallucination in large language models (LLM) by preparing responses in real, verified data rather than relying completely on previously learned knowledge. By retrieving appropriate material from external sources, RAG provides specific, up-to-date context that directly connects to queries. Which reduces the model's bias to previously learned knowledge or rely on obsolete information. This retrieval also permits verifiable source acknowledgement, enabling users to verify facts and make the model accountable. [3], (Tonmoy et al., 2024)

Further, **Agent-based LLM** models can help mitigate hallucination by using agents that verify facts, cross-reference sources, and validate information before finalizing answers. By running real-time checks and focusing on task-specific information, agents can help prevent the model from fabricating with previously learned knowledge. (Greyling, 2024)

Both Agent-based systems and Retrieval-Augmented Generation (RAG) improve referencing and reasoning in LLMs by relating real-time data verification with contextually appropriate knowledge retrieval. (Tonmoy et al., 2024; Han et al., 2024; Greyling, 2024; Stakelum, 2024)

Language models demonstrate success across range of text generation, their **ability to reasoning** is often seen as a limitation and proven that model size scale does not change its reasoning capabilities. (Wang et al., 2023) To improve reasoning in Large Language models, is necessary to build trust on LLMs. Pre-train LLMs often perform better for arithmetic, symbolic and logical reasoning tasks when they are not used as zero-shot reasoners. (Kojima et al., 2023)

On the other hand, some research has shown that **hallucination is inevitable**. (Xu et al., 2024). Further it shows that LLMs will Hallucinate on what they cannot compute. This theoretical understanding means that striving for a completely hallucination-free LLM is an unattainable goal. The focus of research and development should shift towards mitigating and managing hallucinations rather than trying to eliminate them entirely. (Xu et al., 2024)

In areas fields like art, literature, and design **hallucination is expected by the LLM** itself. Hallucinations, in these contexts, can lead to novel perspectives and creative ideas that a purely factual system might miss. When brainstorming or exploring new concepts, a certain level of deviation from established facts can help to generate new ideas. (Huang et al., 2023; Xu et al., 2024)

Methodology

Following methodology has been identified as the methodology for this research,

- **A Complete Literature Review**

Conduct a complete review of existing research on hallucinations, LLMs, RAG, and agent-based AI systems.

- Understanding Hallucination in LLMs
- Existing Methods to Reduce Hallucinations
- Understand Evaluation metrics of Hallucinations
- Agent Architecture
- Retrieval-Augmented Generation (RAG)
- Referencing and Reasoning
- Ethical and Practical Considerations

- **Problem Definition and Dataset Selection**

Define the specific hallucination types to target and select datasets relevant to high-stakes environment.

- Identify hallucination type impacts largely on high-stakes environments
- Consult domain experts for further information
- Dataset validation
- Pilot testing

- **Develop an Agent-Based Framework**

Design and implement agents-based framework for generalized use.

- Outline the architecture for the agent-based framework
- Selection of tools and techniques
- Develop selected framework(s)

- **Integrate Retrieval-Augmented Generation (RAG)**

Implement RAG to support LLMs with access to a reliable external knowledge base.

- **Enhance Referencing and Reasoning**


Develop referencing, attribution, and logic-based reasoning when using agents and RAG

- **Evaluation Metrics and Testing**

Identify and evaluate selected model on high-stakes datasets.

- Quantitative Analysis
 - Factual Accuracy Metrics
 - Hallucination Detection methods
 - Reference Precision
- Qualitative Analysis
 - Expert Review
 - Error Analysis
- User Trust and Interpretability Analysis

Timeline



Expected Outcome

- 1. Reduction in Hallucination Rate:** The primary expected outcome of this research is a considerable reduction in the hallucination or frequency of hallucinations within large language model outputs, particularly when tested in high-stakes environments. The research aims to create a framework that produces more accurate, factually correct, and reliable outputs. This reduction will be validated over traditional LLM outputs.
- 2. Enhanced Referencing and Traceability:** Another outcome of this research is for referencing and traceability. The framework needs to provide accurate references and citations that link responses directly to data sources. This capability will allow users to trace back information. Enhanced referencing ensures that outputs are verifiable, which is essential in high-stakes environments.
- 3. Improved Reasoning Capabilities:** The research aims to boost the LLM’s logic-based reasoning capabilities. The expected result is logically correct answers, addressing queries with confidence and accuracy.

Conclusion

This research aims to find an innovative approach to minimize hallucinations in LLMs through integration of agent-based systems with retrieval-augmented generation (RAG). Furthermore, it looks for ways to enhance referencing and reasoning mechanisms in existing LLMs. The proposed solution framework will reduce hallucinations, supported by verifiable references and improved logical reasoning in high-stakes environments.

References

- Dhuliawala, S., Komeili, M., Xu, J., Raileanu, R., Li, X., Celikyilmaz, A. and Weston, J. 2023. Chain-of-Verification Reduces Hallucination in Large Language Models.
- Greyling, C. 2024. An AI Agent Architecture & Framework Is Emerging. *Medium*. [Online]. [Accessed 27 October 2024]. Available from: <https://cobusgreyling.medium.com/an-ai-agent-architecture-framework-is-emerging-addae3804f23>.
- Han, K., Kuang, K., Zhao, Z., Ye, J. and Wu, F. 2024. Causal Agent based on Large Language Model.
- Huang, L., Yu, W., Ma, W., Zhong, W., Feng, Z., Wang, H., Chen, Q., Peng, W., Feng, X., Qin, B. and Liu, T. 2023. A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions.
- Kojima, T., Gu, S.S., Reid, M., Matsuo, Y. and Iwasawa, Y. 2023. Large Language Models are Zero-Shot Reasoners.
- Manakul, P., Liusie, A. and Gales, M.J.F. 2023. SelfCheckGPT: Zero-Resource Black-Box Hallucination Detection for Generative Large Language Models.
- Stakelum, J.L. 2024. The End of AI Hallucinations: A Big Breakthrough in Accuracy for AI Application Developers. *Medium*. [Online]. [Accessed 27 October 2024]. Available from: <https://medium.com/@JamesStakelum/the-end-of-ai-hallucinations-a-breakthrough-in-accuracy-for-data-engineers-e67be5cc742a>.
- Tonmoy, S.M.T.I., Zaman, S.M.M., Jain, V., Rani, A., Rawte, V., Chadha, A. and Das, A. 2024. A Comprehensive Survey of Hallucination Mitigation Techniques in Large Language Models.
- Wang, X., Wei, J., Schuurmans, D., Le, Q., Chi, E., Narang, S., Chowdhery, A. and Zhou, D. 2023. Self-Consistency Improves Chain of Thought Reasoning in Language Models.
- Xu, Z., Jain, S. and Kankanhalli, M. 2024. Hallucination is Inevitable: An Innate Limitation of Large Language Models.